

Database Management Systems

**Dr. Santhosh Manikonda
Department of CSE
School of Engineering**

Malla Reddy University, Hyderabad

Syllabus

UNIT - V

- Concurrency Control Techniques:
 - Two-Phase Locking Techniques for Concurrency Control
 - Concurrency Control Based on Timestamp Ordering.
- Disk Storage, Basic File Structures:
 - Introduction
 - Secondary Storage Devices
 - Buffering of Blocks
 - Placing File Records on Disk
 - Operations on Files.

Concurrency Control Protocols

- Primary goal of concurrency protocols is to achieve consistency
- This goal is achieved by using different protocols
 - Shared/Exclusive Lock
 - Two – Phase Locking

Shared/Exclusive Lock

- **Shared Lock (S):** In shared lock, a transaction is allowed to **only read**
 - **Shared lock allows another transaction(s) for shared lock only**
- **Exclusive Lock (X):** In exclusive lock, a transaction is allowed to **read and write**
 - **Exclusive lock does not allow any lock until it is released**
 - If a lock cannot be granted, the requesting transaction is made to **wait till** all incompatible locks held by other transactions have been released. The lock is then granted.

Shared/Exclusive Lock

		Request	
		S	X
Grant	S	YES	NO
	X	NO	NO

Shared/Exclusive Lock

T1	T2
X(A)	
R(A)	
W(A)	
U(A)	
	S(A)
	R(A)
	U(A)
X(B)	
R(B)	
W(B)	
U(B)	

Shared/Exclusive Lock

- Problems in Shared Lock/Exclusive Lock
 - May not be sufficient to achieve serializable schedule
 - May not be recoverable
 - May not be free from deadlock
 - May not be free from starvation

Shared/Exclusive Lock

- May not be recoverable

T1	T2
X(A)	
R(A)	
W(A)	
U(A)	
	X(A)
	R(A)
	W(A)
	U(A)
	Commit
X(B)	
R(B)	
W(B)	
U(B)	
Failed	

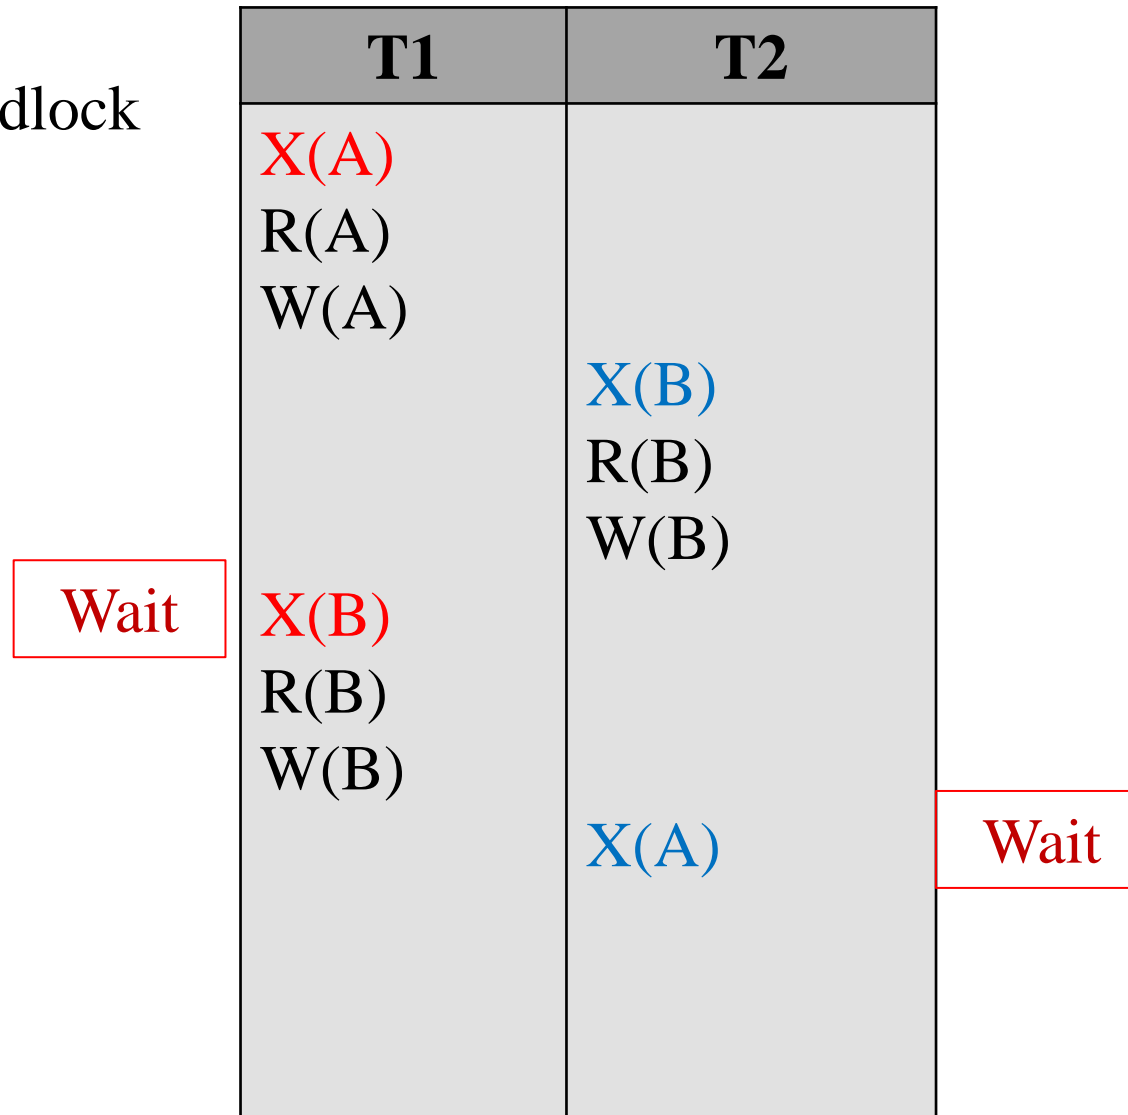
Shared/Exclusive Lock

- May not be recoverable

T1	T2
X(A) R(A) W(A) U(A)	<div>X(A) R(A) W(A) U(A) Commit</div>
X(B) R(B) W(B) U(B) Failed	

Shared/Exclusive Lock

- May not be free from deadlock



Shared/Exclusive Lock

- May not be free from starvation

T1	T2	T3	T4
S(A)	X(A)	S(A)	
U(A)			
		U(A)	S(A)


Two – Phase Lock (2PL)

- Growing Phase:
 - In this phase a transaction may acquire locks
 - Transaction may not release locks
- Shrinking Phase:
 - In this phase Transaction may release locks
 - Transaction may not obtain locks
- The protocol assures serializability.
- It can be proved that the transactions can be serialized in the order of their lock points (i.e., the point where a transaction acquired its final lock).

Two – Phase Lock (2PL) – Lock Conversions

- Two-phase locking with lock conversions:
 - First Phase:
 - can acquire a lock-S on item
 - can acquire a lock-X on item
 - can convert a lock-S to a lock-X (upgrade)
 - Second Phase:
 - can release a lock-S
 - can release a lock-X
 - can convert a lock-X to a lock-S (downgrade)
- This protocol assures serializability. But still relies on the programmer to insert the various locking instructions.

Two – Phase Lock (2PL)



Two – Phase Lock (2PL)

T1	T2
S(A) R(A)	
X(B) R(B) W(B) . . .	S(A) R(A) W(A) S(D) R(D) . . .
U(A) U(B)	

Automatic Acquisition of Locks

- A transaction T_i issues the standard read/write instruction, without explicit locking calls.
- The operation **read(D)** is processed as:

```
    if  $T_i$  has a lock on D
        then
            read(D)
    else
        begin
            if necessary wait until no other
            transaction has a lock-X on D
            grant  $T_i$  a lock-S on D;
            read(D)
        end
```


Automatic Acquisition of Locks

- The operation write(D) is processed as:

```
if Ti has a lock-X on D
    then
        write(D)
    else begin
        if necessary wait until no other transaction
        has any lock on D,
        if Ti has a lock-S on D
            then
                upgrade lock on D to lock-X
            else
                grant Ti a lock-X on D
        write(D)
    end;
```

- All locks are released after commit or abort


Two – Phase Lock (2PL)

- Strict 2PL
 - Strict 2PL should satisfy basic 2PL
 - All **Exclusive Locks** should be held until Commit/Abort
- Rigorous 2PL:
 - Rigorous 2PL should satisfy basic 2PL
 - All **Exclusive Locks, Shared Locks** should be held until Commit/Abort

Implementation of Locking

- A lock manager can be implemented as a **separate process** to which transactions send lock and unlock requests
- The lock manager replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock)
- The requesting transaction waits until its request is answered
- The lock manager maintains a data-structure called a **lock table** to record granted locks and pending requests
- The lock table is usually implemented as an **in-memory hash table** indexed on the name of the data item being locked

Lock Table



- Dark blue rectangles indicate granted locks;
- Light blue indicate waiting requests
- Lock table also records the type of lock granted or requested
- New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- Unlock requests result in the request being deleted, and later requests are checked to see if they can now be granted
- If transaction aborts, all waiting or granted requests of the transaction are deleted
- Lock manager may keep a list of locks held by each transaction, to implement this efficiently

Timestamp-Based Protocol

- Each transaction is issued a timestamp when it enters the system
 - It is a unique value assigned to every transaction
 - It tells the order in which a transaction has entered the system
 - Transaction: T_i
 - Time stamp: $TS(T_i)$
- If an old transaction T_i has time-stamp $TS(T_i)$, a new transaction T_j is assigned time-stamp $TS(T_j)$ such that $TS(T_i) < TS(T_j)$
- The protocol manages concurrent execution such that the time-stamps determine the serializability order – *Older transactions are executed first*

Timestamp-Based Protocol

- In order to assure such behavior, the protocol maintains for each data A two timestamp values:
- W-timestamp(A)
 - **WTS(A)** is the largest time-stamp of any transaction that executed write(A) successfully – **Last transaction which performed Write successfully**
- R-timestamp(A)
 - **RTS(A)** is the largest time-stamp of any transaction that executed read(A) successfully – **Last transaction which performed Read successfully**

Timestamp-Based Protocol

- Time Stamp of Transaction $TS(T_i)$

10:00	10:05	10:07	(Time of Transaction)
T1	T2	T3	T_i
100	120	134	$TS(T_i)$
<i>Oldest</i>		<i>Youngest</i>	

Timestamp-Based Protocol

- Time stamp of Data Item $RTS(A)$

09:00	09:03	09:15	
T1	T2	T3	Ti
10	12	24	TS(T _i)
R(A)	R(A)	<div>R(A)</div>	

$$RTS(A) = 24$$

Timestamp-Based Protocol

- Time stamp of Data Item $WTS(A)$

09:00	09:03	09:15	
T1	T2	T3	Ti
10	12	24	TS(T _i)
W(A)		W(A)	

$$WTS(A) = 12$$

Timestamp-Based Protocol

- Suppose a transaction T_i issues a **Read(A)**
- If $TS(T_i) < WTS(A)$, then T_i needs to read a value of A that was already overwritten.
 - Hence, the read operation is rejected, and T_i is **rolled back**
- If $TS(T_i) \geq WTS(A)$, then the read operation is executed, and set $RTS(A) = \max\{RTS(A), TS(T_i)\}$

Example
 $TS(T_i) < WTS(A)$
 $900 < 903$

09:00	09:03
T_i	T_x
$R(A)$	$W(A)$. . . XXXX

Timestamp-Based Protocol

- Suppose a transaction T_i issues a **Read(A)**
- If $TS(T_i) < WTS(A)$, then T_i needs to read a value of A that was already overwritten.
 - Hence, the read operation is rejected, and T_i is rolled back
- If $TS(T_i) \geq WTS(A)$, then the read operation is **executed**, and set **$RTS(A) = \max\{RTS(A), TS(T_i)\}$**

Example $TS(T_i) > WTS(A)$ 910 > 903
--

09:10	09:03
T _i	T _x
R(A)	W(A)

Timestamp-Based Protocol

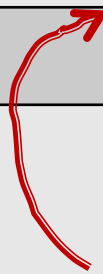
- Suppose a transaction T_i issues a Write(A)
- If $TS(T_i) < RTS(A)$, then the value of A that T_i is producing was needed previously, and the system assumed that that value would never be produced.
 - Hence, the write operation is rejected, and T_i is rolled back
- If $TS(T_i) < WTS(A)$, then T_i is attempting to write an obsolete value of A
 - Hence, the write operation is rejected, and T_i is rolled back
- Otherwise, the write operation is executed, and set $WTS(A) = \max\{WTS(A), TS(T_i)\}$

Timestamp-Based Protocol

- Suppose a transaction T_i issues a **Write(A)**
- If $TS(T_i) < RTS(A)$, then the value of A that T_i is producing was needed previously, and the system assumed that that value would never be produced.
 - Hence, the write operation is rejected, and T_i is **rolled back**
- If $TS(T_i) < WTS(A)$, then T_i is attempting to write an obsolete value of A
 - Hence, the write operation is rejected, and T_i is rolled back
- Otherwise, the write operation is executed, and set $WTS(A) = \max\{WTS(A), TS(T_i)\}$

Example	
$TS(T_i) < RTS(A)$	
900	< 903

09:00	09:03
T_i	T_x
$W(A)$	$R(A)$




Timestamp-Based Protocol

- Suppose a transaction T_i issues a **Write(A)**
- If $TS(T_i) \leq RTS(A)$, then the value of A that T_i is producing was needed previously, and the system assumed that that value would never be produced.
 - Hence, the write operation is rejected, and T_i is rolled back
- If $TS(T_i) < WTS(A)$, then T_i is attempting to write an obsolete value of A
 - Hence, the write operation is rejected, and T_i is **rolled back**
- Otherwise, the write operation is executed, and set $WTS(A) = \max\{WTS(A), TS(T_i)\}$

Example	
$TS(T_i) < WTS(A)$	
900	< 903

09:00	09:03
T _i	T _x
W(A)	W(A)



Timestamp-Based Protocol

- Suppose a transaction T_i issues a **Write(A)**
- If $TS(T_i) \leq RTS(A)$, then the value of A that T_i is producing was needed previously, and the system assumed that that value would never be produced.
 - Hence, the write operation is rejected, and T_i is rolled back
- If $TS(T_i) < WTS(A)$, then T_i is attempting to write an obsolete value of A
 - Hence, the write operation is rejected, and T_i is rolled back
- Otherwise, the write operation is **executed**, and set $WTS(A) = \max\{WTS(A), TS(T_i)\}$

Example
 $TS(T_i) > RTS(A)$

Example
 $TS(T_i) > WTS(A)$

Timestamp-Based Protocol - Properties

- It ensures conflict serializability
- It ensures view serializability
- Free from deadlock
- Possibility of dirty read and irrecoverable schedule